

# Week2\_Capstone\_Project\_Final

June 4, 2020

## 0.1 # Capstone Project: The Battle of Neighborhoods

## San Francisco Police Department Incident Reports: 2018 to Present(01/06/2020) ### Author : Jithin Prakash Kolamkolly \*\*\*

Police Department is collating data about the criminal incidents happening in and around San Francisco, that data is published in 'The office of the chief Data Officer – City and County of San Francisco' (<https://data.sfgov.org/Public-Safety/Police-Department-Incident-Reports-2018-to-Present/wg3w-h783>) website. This data is directly pulled from the website to do modelling and to analyze the data

CRISP DM methodology is used in analysis and prediction

- Collecting Data
- Explore and Understand Data
- Data Preparation and pre-processing
- Modelling
- Evaluation and testing

## 0.2 ##### Importing Libraries

```
[0]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

## 1 1. Collect Data

Importing Data from website (<https://data.sfgov.org/api/views/wg3w-h783/rows.csv>)

```
[0]: !wget -q -O 'Police_incidents.csv' https://data.sfgov.org/api/views/wg3w-h783/
    →rows.csv
print('Data fetched from website')
```

Data fetched from website

---

## 2 2. Explore and Understand Data

Reading CSV to Data Frame and performing operations

```
[0]: #Reading from the CSV file to the data frame
df_Police = pd.read_csv('Police_incidents.csv')
#Printing the shape of the Raw Data
print('\nRaw data has %d Rows and %d Columns\n'% df_Police.shape)
#Displaying the first 5 rows
df_Police.head()
```

Raw data has 351980 Rows and 36 Columns

```
[0]:      Incident Datetime ... Areas of Vulnerability, 2016
0  2019/05/01 01:00:00 AM ...                1.0
1  2019/06/22 07:45:00 AM ...                2.0
2  2019/06/03 04:16:00 PM ...                2.0
3  2018/11/16 04:34:00 PM ...                2.0
4  2019/05/27 02:25:00 AM ...                1.0
```

[5 rows x 36 columns]

```
[0]: #Check for Duplicate Entries
df_Police.duplicated().sum()
```

```
[0]: 0
```

```
[0]: #Check the details using describe method
df_Police.describe()
```

```
[0]:      Incident Year ... Areas of Vulnerability, 2016
count  351980.000000 ...                333307.000000
mean    2018.703597 ...                1.549637
std       0.699340 ...                0.497531
min     2018.000000 ...                1.000000
25%     2018.000000 ...                1.000000
50%     2019.000000 ...                2.000000
75%     2019.000000 ...                2.000000
max     2020.000000 ...                2.000000
```

[8 rows x 20 columns]

### 3 3. Data Preparation and Pre-processing

Refer the documentation for the Column Header / Feature / Attribute description

```
[0]: #Dropping the data that is not needed for the analysis
df_Police.drop(['Report Datetime', 'Row ID', 'Incident ID', 'Incident Number',
```

```

        'CAD Number','Report Type Code','Report Type',
→Description','Filed Online','Incident Code',
        'Incident Subcategory','CNN','Supervisor District','point','SF',
→Find Neighborhoods','Current Police Districts',
        'Current Supervisor Districts','Analysis Neighborhoods','HSOC',
→Zones as of 2018-06-05','OWED Public Spaces',
        'Central Market/Tenderloin Boundary Polygon - Updated','Parks',
→Alliance CPSI (27+TL sites)',
        'ESNCAG - Boundary File','Areas of Vulnerability',
→2016'],axis=1,inplace=True)

#Printing the shape of the Processed Data
print('\nColumn Drop - Data has %d Rows and %d Columns'% df_Police.shape)

#Dropping the Rows that has NaN or Null values
df_Police.dropna(subset = ['Incident Datetime','Incident Date','Incident',
→Time','Incident Year','Incident Day of Week',
        'Incident Category','Incident',
→Description','Resolution','Intersection',
        'Police District','Analysis',
→Neighborhood','Latitude','Longitude'],inplace=True,axis=0)

# Adding Month Column on to the DataFrame
df_Police.insert(3,'Incident Month',df_Police['Incident Date'].apply(lambda x:
→str(x.split("/")[1])))

#Printing the shape of the Processed Data
print('\nRow Drop - Data has %d Rows and %d Columns\n'% df_Police.shape)

df_Police.head()

```

Column Drop - Data has 351980 Rows and 13 Columns

Row Drop - Data has 333185 Rows and 14 Columns

```

[0]:      Incident Datetime Incident Date  ...  Latitude  Longitude
0  2019/05/01 01:00:00 AM    2019/05/01  ...   37.762569 -122.499627
1  2019/06/22 07:45:00 AM    2019/06/22  ...   37.780535 -122.408161
2  2019/06/03 04:16:00 PM    2019/06/03  ...   37.721600 -122.390745
3  2018/11/16 04:34:00 PM    2018/11/16  ...   37.794860 -122.404876
4  2019/05/27 02:25:00 AM    2019/05/27  ...   37.797716 -122.430559

```

[5 rows x 14 columns]

### Crimes per Category

Lets check the number of crimes by category

```
[0]: df_category = pd.DataFrame(df_Police['Incident Category'].value_counts())
df_category=df_category.reset_index().rename(columns={'index' : 'Incident_
→Category','Incident Category':'Incident Count'})
df_category.head(15)
```

```
[0]:
```

|    | Incident Category                        | Incident Count |
|----|--|----------------|
| 0  | Larceny Theft                            | 100845         |
| 1  | Other Miscellaneous                      | 26326          |
| 2  | Non-Criminal                             | 21453          |
| 3  | Assault                                  | 20712          |
| 4  | Malicious Mischief                       | 20516          |
| 5  | Burglary                                 | 16200          |
| 6  | Motor Vehicle Theft                      | 13056          |
| 7  | Warrant                                  | 12387          |
| 8  | Fraud                                    | 10246          |
| 9  | Lost Property                            | 9955           |
| 10 | Drug Offense                             | 8792           |
| 11 | Robbery                                  | 8244           |
| 12 | Missing Person                           | 8009           |
| 13 | Recovered Vehicle                        | 7753           |
| 14 | Offences Against The Family And Children | 6815           |

Plotting the data to see the frequency of all crime categories


```
[0]: #df_category.sort_values('Incident Category',ascending=True,inplace=True)

ax = df_category.plot(kind='barh',figsize=(10,20),color='steelblue')
ax.set_title('Frequency of Crimes per Categoty',fontsize=16)
ax.set_yticklabels( df_category['Incident Category'],fontsize=10)
ax.set_ylabel('Frequency')
ax.tick_params(labelsize=12)

ax.invert_yaxis()
ax.xaxis.tick_top()

ax.xaxis.grid(True)
ax.get_legend().remove()

#plt.savefig('plt_1.png',bbox_inches='tight')
```




**Time Crime Incident - Year/Month/Time/Day analysis    Month**

```
[0]: import calendar

df_month = pd.DataFrame(df_Police['Incident Month'].value_counts()).
    →reset_index()
df_month.sort_values('index',ascending=True,inplace=True)

ax = df_month.plot(kind='bar',figsize=(10,5))
ax.set_title('Frequency of Crimes per Month',fontsize=16)
ax.set_xticklabels( df_month['index'].apply(lambda x: calendar.
    →month_abbr[int(x)]),fontsize=10)
ax.set_ylabel('Frequency')
ax.tick_params(labelsize=10)

ax.yaxis.grid(True)
ax.get_legend().remove()
```



### Day

```
[0]: day_seq = ['Sunday', 'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
    →'Saturday']
df_day = pd.DataFrame(df_Police['Incident Day of Week'].value_counts()).
    →reset_index()
df_day.sort_values('index',ascending=True,inplace=True)
df_day['index'] = pd.Categorical(df_day['index'], categories=day_seq,
    →ordered=True)
df_day.sort_values('index',inplace=True)


ax = df_day.plot(kind='bar',figsize=(10,5),width=.35)
```

```

ax.set_title('Frequency of Crimes per weekday',fontsize=16)
ax.set_xticklabels( df_day['index'],fontsize=10)
ax.set_ylabel('Frequency')
ax.tick_params(labelsize=10)

ax.yaxis.grid(True)
ax.get_legend().remove()

```




### Time of the Day

```

[0]: df_time = pd.DataFrame(pd.to_datetime(df_Police['Incident Time'],_
    →errors='coerce').dt.hour)
df_time = pd.DataFrame(df_time['Incident Time'].value_counts()).reset_index().
    →sort_values('index')

ax = df_time.plot(kind='bar',figsize=(12,5),color='g',width = 1.2)
ax.set_title('Frequency of Crimes per Time',fontsize=16)
ax.set_xticklabels( df_time['index'],fontsize=10)
ax.tick_params(labelsize=10)
ax.set_xlabel('Hour of the Day',fontsize=10)
ax.set_ylabel('Frequency',fontsize=10)
ax.set_xlim(0)
ax.yaxis.grid(True)
ax.get_legend().remove()

```



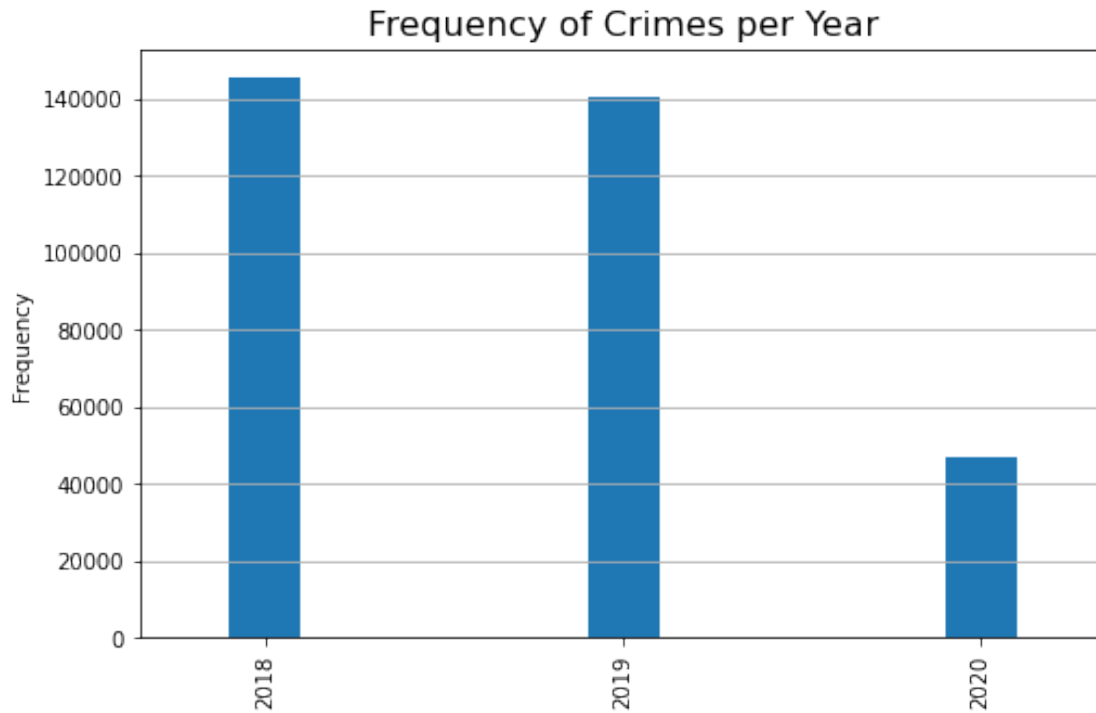
### Year

```
[0]: df_year = pd.DataFrame(df_Police['Incident Year'].value_counts())

ax = df_year.plot(kind='bar',figsize=(8,5),width=0.2,align='center')
ax.set_title('Frequency of Crimes per Year',fontsize=16)
ax.set_xticklabels( df_year.index,fontsize=10)
#ax.set_xlabel('Year',fontsize=10)
ax.set_ylabel('Frequency')
ax.tick_params(labelsize=10)

ax.yaxis.grid(True)
ax.get_legend().remove()
```






### Crime Frequency on Different WeekDays

```
[0]: df_time_Day = df_Police[['Incident Time', 'Incident Day of Week']]
df_time = pd.DataFrame(pd.to_datetime(df_Police['Incident Time'],
    →errors='coerce').dt.hour).rename(columns={'Incident Time': 'TimeH'})
df_time_Day = pd.concat([df_time_Day, df_time], axis=1)

[0]: df_time_Day_group = df_time_Day.groupby(['TimeH', 'Incident Day of Week']).
    →count()['Incident Time'].unstack()
fig, ax = plt.subplots(figsize=(15, 7))
df_time_Day_group.plot(ax=ax)
ax.set_title('Hourly Crime Incident frequency on each Week Day')
ax.set_xlabel('Hour - Day')
ax.set_ylabel('Frequency')
ax.set_xlim(0)
ax.grid()
```




### Crime per Police District

```
[0]: df_district = pd.DataFrame(df_Police['Police District'].value_counts()).
      ↪reset_index()
df_district.sort_values('index',inplace=True)
df_district

ax = df_district.plot(kind='bar',figsize=(10,5),width=0.4)
ax.set_title('Frequency of Crimes per Police District',fontsize=16)
ax.set_xticklabels( df_district['index'],fontsize=10)
ax.set_xlabel('Police District',fontsize=10)
ax.set_ylabel('Frequency')
ax.tick_params(labelsize=10)

ax.yaxis.grid(True)
ax.get_legend().remove()
```



### 3.0.1 Populate a word-cloud to understand the most occurring crime types

```
[50]: #Install Word-Cloud if not available
!conda install -c conda-forge wordcloud --yes

from wordcloud import WordCloud
from wordcloud import STOPWORDS


stopWords = set(STOPWORDS)
stopWords.add('Found')

wc = WordCloud(background_color = 'white', max_words=20000, width = 1000,
               height = 700, stopwords=stopWords)
wc.generate(str(df_Police['Incident Description']))
plt.figure(figsize=(15,8))
plt.imshow(wc,interpolation='bilinear')
plt.axis('off')
plt.title('Description of the Crime\n', fontsize = 15)

plt.show()
```

/bin/bash: conda: command not found

## Description of the Crime



## 4 4. Modelling

### Pre-Processing

Filtering relevant features to X and Y arrays

```
[0]: X = df_Police[['Incident Datetime', 'Police District', 'Latitude', 'Longitude']].  
      ↪values  
      y = df_Police['Incident Category'].values  
      print(X[:5])  
      print(y[:5])
```

```
[['2019/05/01 01:00:00 AM' 'Taraval' 37.76256939715695  
 -122.49962745519908]  
 ['2019/06/22 07:45:00 AM' 'Southern' 37.7805353858225  
 -122.40816079455212]  
 ['2019/06/03 04:16:00 PM' 'Bayview' 37.72159985216247  
 -122.39074534279013]  
 ['2018/11/16 04:34:00 PM' 'Central' 37.79485953222834
```

```
-122.40487561154785]
['2019/05/27 02:25:00 AM' 'Northern' 37.79771621229674
-122.43055896140595]]
['Offences Against The Family And Children' 'Non-Criminal'
'Missing Person' 'Offences Against The Family And Children' 'Assault']
```

### Pre processing the Data - Features in X

```
[0]: from sklearn.preprocessing import LabelEncoder

Label_Encoder = LabelEncoder()
X[:,0]=Label_Encoder.fit_transform(X[:,0])
Label_Encoder.fit(['Bayview', 'Central', 'Ingleside', 'Mission', 'Northern', 'Out of_
→SF', 'Park', 'Richmond', 'Southern', 'Taraval', 'Tenderloin'])
X[:,1]=Label_Encoder.fit_transform(X[:,1])
```

### Splitting Test and Train Data

```
[0]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
→random_state=3)
print ('Training Data size for X :{}\nTesting Data size for X :{}\n'.format(
→X_train.shape,X_test.shape))
print ('Training Data size for y :{}\nTesting Data size for y :{}\n'.format(
→y_train.shape,y_test.shape))
```

Training Data size for X :(233229, 4)

Testing Data size for X :(99956, 4)

Training Data size for y :(233229,)

Testing Data size for y :(99956,)

### K-Nearest Neighbour \_\_\_\_

```
[0]: #Finding K
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from sklearn.metrics import jaccard_similarity_score

Ks=10
mean_acc = np.zeros((Ks-1))
std_acc = np.zeros((Ks-1))

for n in range(1,Ks):
    #Train Model and Predict
    neigh = KNeighborsClassifier(n_neighbors = n).fit(X_train,y_train)
    yhat=neigh.predict(X_test)
    mean_acc[n-1] = metrics.accuracy_score(y_test, yhat)
    std_acc[n-1]=np.std(yhat==y_test)/np.sqrt(yhat.shape[0])
```

```
mean_acc
print( "The best accuracy was with", mean_acc.max(), "with k=", mean_acc.
      ↳argmax()+1)
```

The best accuracy was with 0.24718876305574453 with k= 9

It is found that the accuracy is higher = 0.25 when k is 9

Prediction

```
[0]: neigh = KNeighborsClassifier(n_neighbors = 9).fit(X_train,y_train)
      yhat_KNN = neigh.predict(X_test)

[0]: print("Train set Accuracy : ", metrics.accuracy_score(y_train, neigh.
      ↳predict(X_train)))
      print("Test set Accuracy : ", metrics.accuracy_score(y_test, yhat_KNN))
      print("F1 Accuracy : ", metrics.f1_score(y_test, yhat_KNN, average='weighted'))
      print("Jaccard Similarity Score : ", jaccard_similarity_score(y_test, yhat_KNN))
```

```
Train set Accuracy : 0.354685738051443
Test set Accuracy : 0.24718876305574453
F1 Accuracy : 0.17138718298316208
Jaccard Similarity Score : 0.24718876305574453
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:664:
FutureWarning: jaccard_similarity_score has been deprecated and replaced with
jaccard_score. It will be removed in version 0.23. This implementation has
surprising behavior for binary and multiclass classification tasks.
FutureWarning)
```

**Decision Tree \*\*\***

```
[0]: from sklearn.tree import DecisionTreeClassifier
      cat_Tree = DecisionTreeClassifier(criterion="gini", max_depth = 80)

      cat_Tree.fit(X_train,y_train)

[0]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                             max_depth=80, max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort='deprecated',
                             random_state=None, splitter='best')
```

Prediction

```
[0]: pred_Tree = cat_Tree.predict(X_test)
      print("DecisionTrees's Accuracy: ", metrics.accuracy_score(y_test, pred_Tree))
      print("Jaccard Similarity Score : ", jaccard_similarity_score(y_test,
      ↳pred_Tree))
```

DecisionTrees's Accuracy: 0.24465764936572093  
Jaccard Similarity Score : 0.24465764936572093

/usr/local/lib/python3.6/dist-packages/sklearn/metrics/\_classification.py:664:  
FutureWarning: jaccard\_similarity\_score has been deprecated and replaced with  
jaccard\_score. It will be removed in version 0.23. This implementation has  
surprising behavior for binary and multiclass classification tasks.  
FutureWarning)

```
[0]: #X = df_Police[.values

'''
from sklearn.externals.six import StringIO
import pydotplus
import matplotlib.image as mpimg
from sklearn import tree
%matplotlib inline

dot_data = StringIO()
filename = "cat_tree.png"
featureNames = ['Incident Datetime', 'Police District', 'Latitude', 'Longitude']
targetNames = df_Police["Incident Category"].unique().tolist()
out=tree.export_graphviz(cat_Tree, feature_names=featureNames,
    →out_file=dot_data, class_names= np.unique(y_train), filled=True,
                        special_characters=True, rotate=False)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png(filename)
img = mpimg.imread(filename)
plt.figure(figsize=(100, 200))
plt.imshow(img, interpolation='nearest')

'''
```

```
[0]: '\nfrom sklearn.externals.six import StringIO\nimport pydotplus\nimport
matplotlib.image as mpimg\nfrom sklearn import tree\n%matplotlib inline
\n\ndot_data = StringIO()\nfilename = "cat_tree.png"\nfeatureNames = ['Incident
Datetime', '\nPolice District', '\nLatitude', '\nLongitude']\ntargetNames =
df_Police["Incident Category"].unique().tolist()\nout=tree.export_graphviz(cat_T
ree, feature_names=featureNames, out_file=dot_data, class_names=
np.unique(y_train), filled=True, \n
special_characters=True, rotate=False) \ngraph =
pydotplus.graph_from_dot_data(dot_data.getvalue())
\ngraph.write_png(filename)\nimg =
mpimg.imread(filename)\nplt.figure(figsize=(100,
200))\nplt.imshow(img, interpolation='\nnearest')\n\n'
```

---

**Logistic Regression \*\*\***

```
[0]: print(X_train.shape,y_train.shape)
      print(X_test.shape,y_test.shape)
```

```
(233229, 4) (233229,)
(99956, 4) (99956,)
```

```
[0]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import log_loss
      from sklearn.metrics import jaccard_similarity_score
      from sklearn.metrics import f1_score
```

```
[0]: LR = LogisticRegression(C=0.01, solver='sag').fit(X_train,y_train)
      yhat_LR = LR.predict(X_test)
      yhat_prob = LR.predict_proba(X_test)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_sag.py:330:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  "the coef_ did not converge", ConvergenceWarning)
```

```
[0]: print("Train set Accuracy : ", metrics.accuracy_score(y_train, LR.
      →predict(X_train)))
      print("Test set Accuracy : ", metrics.accuracy_score(y_test, yhat_LR))
      print("Regression F1 Accuracy : ", metrics.f1_score(y_test, yhat_LR,
      →average='weighted'))
      print("Log Loss : ", log_loss(y_test, yhat_prob))
      print("Jaccard Similarity Score : ", jaccard_similarity_score(y_test, yhat_LR))
```

```
Train set Accuracy : 0.3032513109433218
Test set Accuracy : 0.301312577534115
Regression F1 Accuracy : 0.13953491412845728
Log Loss : 2.817651886914775
Jaccard Similarity Score : 0.301312577534115
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:664:
FutureWarning: jaccard_similarity_score has been deprecated and replaced with
jaccard_score. It will be removed in version 0.23. This implementation has
surprising behavior for binary and multiclass classification tasks.
  FutureWarning)
```

### Support Vector Machine \*\*\*

```
[0]: # from sklearn import svm
      # clf=svm.SVC(kernel='rbf')
      # clf.fit(X_train,y_train)
      # yhat_sum = clf.predict(X_test)
```



```
[0]: # print("Train set Accuracy : ", metrics.accuracy_score(y_train, clf.
      ↳predict(X_train)))
# print("Test set Accuracy : ", metrics.accuracy_score(y_test, yhat_sum))
# print("Regression F1 Accuracy : ", metrics.f1_score(y_test, yhat_sum,
      ↳average='weighted'))
# print("Jaccard Similarity Score : ", jaccard_similarity_score(y_test,
      ↳yhat_sum))
```

---

## 4.1 Plot the points on Map

---

In order to plot the incident map in the neighbourhood, we can make use of the latitude longitude data

```
[0]: #!conda install -c conda-forge geopy --yes # uncomment this line if you haven't
      ↳completed the Foursquare API lab
from geopy.geocoders import Nominatim
#!conda install -c conda-forge folium=0.5.0 --yes # uncomment this line if you
      ↳haven't completed the Foursquare API lab
import folium

[0]: df_Police_Short = df_Police.head(200)

[0]: address = 'San Francisco'

geolocator = Nominatim(user_agent="ny_explorer")
location = geolocator.geocode(address)
latitude = location.latitude
longitude = location.longitude
print('The geograpical coordinate of San Francisco City are {}, {}.'.
      ↳format(latitude, longitude))
```

The geograpical coordinate of San Francisco City are 37.7790262, -122.4199061.

```
[36]: # add markers to map
# instantiate a feature group for the incidents in the dataframe
map_SF = folium.Map(location=[latitude, longitude], zoom_start=12)

incidents = folium.map.FeatureGroup()

# loop through the 100 crimes and add each to the incidents feature group
for lat, lng, in zip(df_Police_Short.Latitude, df_Police_Short.Longitude):
    incidents.add_child(
        folium.CircleMarker(
            [lat, lng],
            radius=5, # define how big you want the circle markers to be
```

```

        color='yellow',
        fill=True,
        fill_color='blue',
        fill_opacity=0.6
    )
)

# add pop-up text to each marker on the map
latitudes = list(df_Police_Short.Latitude)
longitudes = list(df_Police_Short.Longitude)
labels = list(df_Police_Short['Incident Category'])

for lat, lng, label in zip(latitudes, longitudes, labels):
    folium.Marker([lat, lng], popup=label).add_to(map_SF)

# add incidents to map
map_SF.add_child(incidents)

```

[36]: <folium.folium.Map at 0x7fbae9937cf8>

```

[37]: from folium import plugins

# let's start again with a clean copy of the map of San Francisco
map_SF = folium.Map(location = [latitude, longitude], zoom_start = 12)

# instantiate a mark cluster object for the incidents in the dataframe
incidents = plugins.MarkerCluster().add_to(map_SF)

# loop through the dataframe and add each data point to the mark cluster
for lat, lng, label, in zip(df_Police_Short.Latitude, df_Police_Short.
    ↳ Longitude, df_Police_Short['Incident Category']):
    folium.Marker(
        location=[lat, lng],
        icon=None,
        popup=label,
    ).add_to(incidents)

# display map
map_SF

```

[37]: <folium.folium.Map at 0x7fbae9904668>

## 4.2 -----

## 5 Using Fourquare to visualize businesses venues

We will make calls to the Foursquare API for different purposes. You will construct a URL to send a request to the API to search for a specific type of venues, to explore a particular business

venue, to explore a Foursquare user, to explore a geographical location, and to get trending venues around a location. Also, you will learn how to use the visualization library, Folium, to visualize the results.

```
[38]: from geopy.geocoders import Nominatim # module to convert an address into
      ↪ latitude and longitude values
import requests # library to handle requests
import random # library for random number generation

# libraries for displaying images
from IPython.display import Image
from IPython.core.display import HTML

# transforming json file into a pandas dataframe library
from pandas.io.json import json_normalize

#!conda install -c conda-forge folium=0.5.0 --yes
import folium # plotting library

print('Folium installed')
print('Libraries imported.')
```

Folium installed  
Libraries imported.

```
[0]: CLIENT_ID = 'INHYLUFFXJR2LZILTPSQJYV4JQYNEQUFZKNWMQ10CWHAEWJR' # your
      ↪ Foursquare ID
CLIENT_SECRET = 'RH2ADKJBT5GL21C1UV5H44HRC3QTQPWBZFYIU5KODVDOA24M' # your
      ↪ Foursquare Secret
VERSION = '20180605'
LIMIT=100
```

We will use the same data frame that we used to plot the geo-location data of crimes. For an example, let's find the Crime Incident of theft, which is

```
[40]: df_Theft= df_Police_Short.loc[df_Police_Short['Incident Category'].str.
      ↪ contains('Vehicle Theft')]
df_Theft.head(3)
```

```
[40]:      Incident Datetime Incident Date  ...  Latitude  Longitude
47    2019/08/21 02:00:00 PM    2019/08/21  ...    37.769007 -122.438338
99    2018/11/11 09:20:00 AM    2018/11/11  ...    37.779459 -122.402377
117   2020/05/23 06:30:00 PM    2020/05/23  ...    37.807483 -122.413975
```

[3 rows x 14 columns]

### Explore the Area of Theft and venues nearby

Let's assume that this location is a theft prone area as per the analysis. Let's try to explore the area and find any venues within 100m radius

```

[0]: neighborhood_latitude = df_Theft.iloc[2][12]
neighborhood_longitude = df_Theft.iloc[2][13]
#Choesn Central Neighbourhood

[42]: radius = 100 # define radius
url = 'https://api.foursquare.com/v2/venues/explore?
    ↳&client_id={}&client_secret={}&v={}&ll={},{}&radius={}&limit={}'.format(
        CLIENT_ID,
        CLIENT_SECRET,
        VERSION,
        neighborhood_latitude,
        neighborhood_longitude,
        radius,
        LIMIT)
url # display URL

[42]: 'https://api.foursquare.com/v2/venues/explore?&client_id=INHYLUFFXJR2LZILTPSQJYV
4JQYNEQUFZKNWMQ10CWHAEWJR&client_secret=RH2ADKJBT5GL21C1UV5H44HRC3QTQPWBZFYIU5K0
DVDOA24M&v=20180605&ll=37.80748251193778,-122.41397500878729&radius=100&limit=10
0'

[0]: results = requests.get(url).json()
#results

[0]: def get_category_type(row):
    try:
        categories_list = row['categories']
    except:
        categories_list = row['venue.categories']

    if len(categories_list) == 0:
        return None
    else:
        return categories_list[0]['name']

[45]: venues = results['response']['groups'][0]['items']
nearby_venues = json_normalize(venues) # flatten JSON

# filter columns
filtered_columns = ['venue.name', 'venue.categories', 'venue.location.lat', '
    ↳venue.location.lng']
nearby_venues = nearby_venues.loc[:, filtered_columns]

# filter the category for each row
nearby_venues['venue.categories'] = nearby_venues.apply(get_category_type,
    ↳axis=1)

# clean columns
nearby_venues.columns = [col.split(".")[1] for col in nearby_venues.columns]

```

```
nearby_venues.head()
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: FutureWarning:
pandas.io.json.json_normalize is deprecated, use pandas.json_normalize instead
```

```
[45]:
```

|   | name                       | categories          | lat       | lng         |
|---|----------------------------|---------------------|-----------|-------------|
| 0 | Hot Spud                   | Restaurant          | 37.807800 | -122.413997 |
| 1 | Big Bus Tours              | Tour Provider       | 37.808323 | -122.414126 |
| 2 | Hotel Zephyr San Francisco | Hotel               | 37.807763 | -122.413222 |
| 3 | Tower Tours San Francisco  | Tour Provider       | 37.807532 | -122.413749 |
| 4 | Alamo Rent A Car           | Rental Car Location | 37.807722 | -122.414738 |

```
[48]: venues_map = folium.Map(location=[neighborhood_latitude,
    ↳ neighborhood_longitude], zoom_start=20) # generate map centred around the
    ↳ Hotel

# add a red circle marker to represent the Hotel
folium.CircleMarker(
    [neighborhood_latitude, neighborhood_longitude],
    radius=10,
    color='red',
    popup='Crime Incident',
    fill = True,
    fill_color = 'red',
    fill_opacity = 0.6
).add_to(venues_map)

# add the Italian restaurants as blue circle markers
for lat, lng, label in zip(nearby_venues.lat, nearby_venues.lng, nearby_venues.
    ↳ categories):
    folium.CircleMarker(
        [lat, lng],
        radius=5,
        color='blue',
        popup=label,
        fill = True,
        fill_color='blue',
        fill_opacity=0.6
    ).add_to(venues_map)

# display map
venues_map
```

```
[48]: <folium.folium.Map at 0x7fbae93e6e10>
```